

# AI Engineer Take-Home Test

## Objective:

Build a prototype system for **document categorisation** and **content extraction (OCR)** using an LLM / VLM, with a focus on **fraud prevention**. The solution should help determine the most relevant category for the file, and extract content from it.

---

## Scenario

We receive diverse unstructured files from customers. Your task is to process these files to:

1. **Categorise** them into meaningful predefined categories.
2. **Extract content** (e.g., key entities, dates, relevant metadata).

We will provide **API\_KEY** for LLM models [via [OpenRouter](#)] with \$100 credits so that you can use them for your experiments. We will also share several example files to give you an idea of what to focus on.

**File categories:** *invoice, marketplace\_listing\_screenshot, chat\_screenshot, website\_screenshot, other.*

---

## Requirements

### 1. Core functionality

- Accept multiple file types (at minimum: PDF, PNG/JPEG).
- Use an LLM / VLM to:
  - Assign a category (either from a predefined set or dynamically inferred).
  - Extract key structured attributes.
- Provide a standardised, structured output format (e.g., JSON) across different files.
- Provide an interface for uploading files and displaying results (we recommend using a framework such as [Streamlit](#))

### 2. Non-functional concerns

- **Upload speed:** minimize latency from file selection to file processing..
  - **Processing speed:** minimise latency to result availability after upload.
  - **Robustness:** demonstrate how you defend against both user errors and the probabilistic nature of the system.
-

## Implementation guidelines

- Use **basic, readable code** (clarity over complexity).
  - Clearly separate **data ingestion, preprocessing, model calls, and postprocessing**.
  - You may use open-source libraries for file handling and model interaction.
  - We also strongly encourage the use of development tools such as Cursor, Copilot, and others to quickly prototype a working solution.
- 

## Deliverables

- A Git repository containing:
    - The source code.
    - A README explaining:
      - How to run the code.
      - Assumptions and design decisions.
      - Known limitations and possible improvements.
  - Example input files (we will provide several, but you may add your own for testing).
  - Example output files in JSON format.
- 

## Evaluation criteria

We will assess (ordered by priority):

1. **Correctness** – Does the solution categorise and extract features as intended?
  2. **Code quality** – Is it clean, modular, and easy to follow?
  3. **Performance awareness** – Are speed and stability considered?
  4. **Guardrails** – Does the solution fail gracefully?
  5. **Extensibility** – Is there a clear path for future improvements?
- 

## Time expectations

Please don't spend more than 3 hours on this. We are not looking for a production-ready system — focus on demonstrating your **problem-solving approach** and ability to work with LLM/VLM tools.